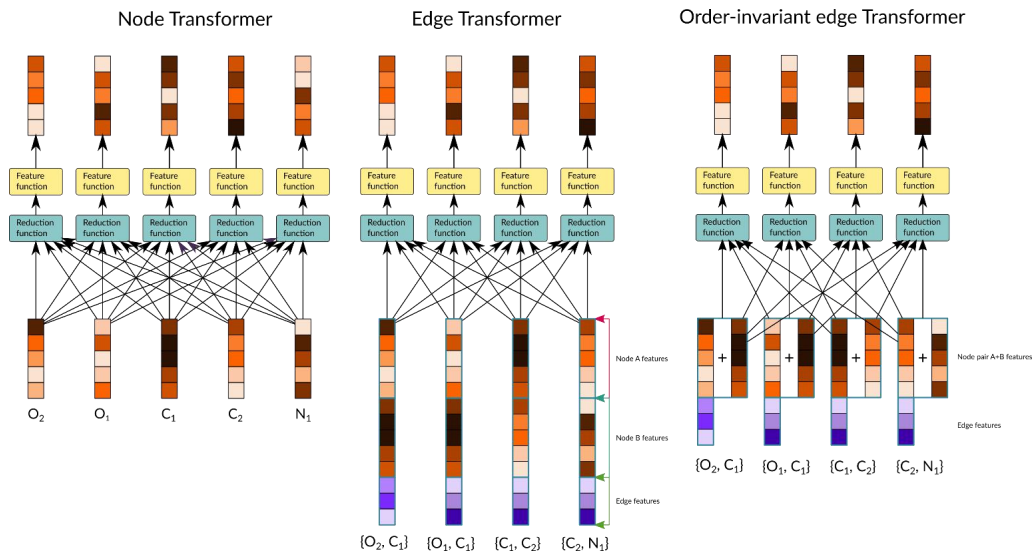
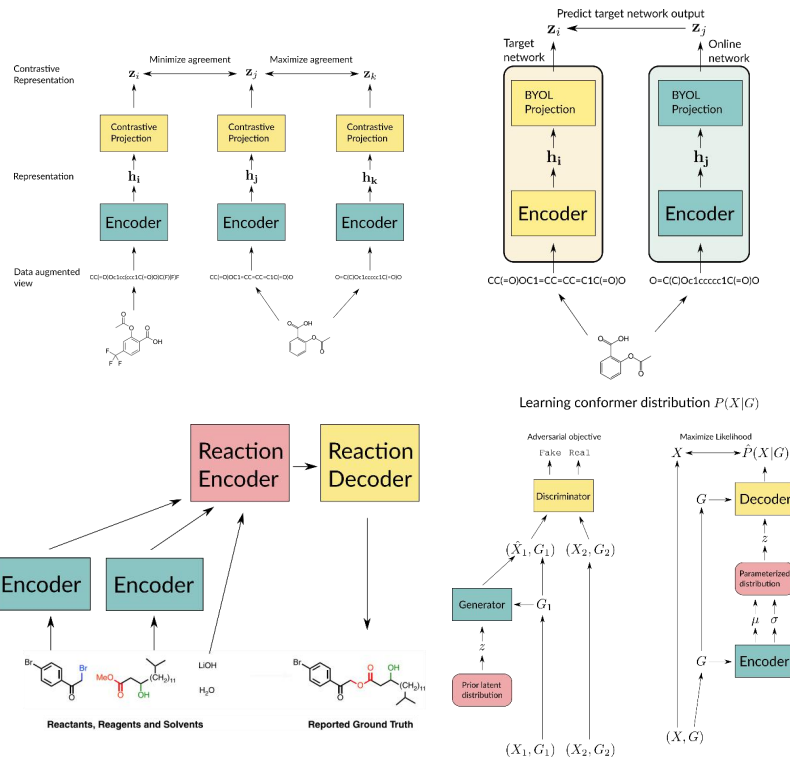


# Optimizing deep learning

# My research



Graph Transformers



Pretraining of  
Molecule encoders

RL  
SE

# The EuroCC National EuroHPC Competence Center Sweden (ENCCS)



The EuroCC National EuroHPC Competence Center Sweden (ENCCS) is a joint initiative between the ten main Swedish research universities and RISE Research Institutes of Sweden. The center is hosted by Uppsala University (UU) on behalf of the consortium and will include the relevant competences at the other nodes. The initiative is funded by the EuroHPC JU, Swedish Research Council (Vetenskapsrådet) and the Swedish Innovation Agency (Vinnova). It is designed to prioritize support based on

- ▶ Needs of academic users with large scale allocations such as PRACE allocations in EU and SNIC allocations in Sweden
- ▶ The current industrial usage of HPC and their future HPC and Artificial Intelligence (AI) needs
- ▶ Needs for training and support to enable a wide range of Swedish users to use the new hardware deployed in pre-exascale systems, in particular Euro HPC JU (pre)exascale system LUMI

# Profilers

- nvprof - deprecated profiler for CUDA
  - Nsight Systems should be used instead
- py-spy - sampling profiler for python code
- line\_profiler - profile each line of a function

# Logging GPU usage

- `nvidia-smi --query-gpu=timestamp,pci.bus_id,utilization.gpu,utilization.memory --format=csv -l 1`
- Small script for online logging: <https://github.com/eryl/gpulog>

# CUDA semantics in PyTorch

- Cuda calls from our frameworks are generally asynchronous
- This might make profiling information at the python-level misleading, the time spent at one location in python code might be due to waiting for a previous call to finish

# NVprof (deprecated)

- Used to profile CUDA API calls
- As a deep learning researcher, this is typically below the level of abstraction we work at
- I've mainly used it to understand if memory copy from host to device is an issue

# Installing nvprof

- If using ubuntu package manager to install NVIDIA software, available as nvidia-profiler
  - apt install nvidia-profiler
- By default requires privileged access
- On the AIDA DGX-2, it seems to work out of the box



# profiling with nvprof

- For the use case of analyzing memory copies, default invocation works fine
- `nvprof --log-file nvprof_log.txt python script.py`
- This will dump profiling output to the text file `nvprof_log.txt`

# nvprof output

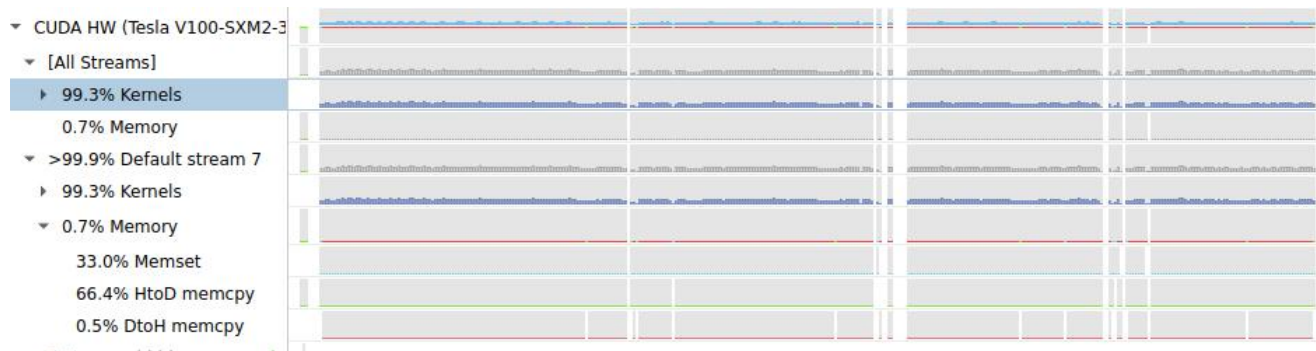
- The output shows what cuda kernels have been run on the GPU(s) how much % of time they have used
- Here I typically look to [CUDA memcpy HtoD], how much time was spent copying data from the host to the device
- For deep learning programs, this is typically quite low (0.12% in this case)

```
==40242== Profiling result:
```

GPU activities:	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
		11.36%	19.0811s	1180113	16.168us	736ns	9.2059ms	void at::native::ve
		9.08%	15.2482s	57720	264.18us	221.13us	6.3191ms	void cudnn::cnn::wg
		6.57%	11.0269s	39000	282.74us	168.87us	7.4494ms	void cudnn::detail:
		6.28%	10.5512s	61390	171.87us	151.21us	7.4208ms	maxwell_scudnn_wino
		5.66%	9.50446s	265156	35.844us	1.6640us	9.3531ms	_ZN2at6native29vect
		4.95%	8.31429s	36136	230.08us	110.79us	8.1986ms	void precomputed_co
		4.52%	7.59856s	111540	68.124us	29.346us	5.8234ms	void cudnn::bn_bw_1
		4.51%	7.57257s	1092780	6.9290us	736ns	6.4271ms	void at::native::ve
		4.08%	6.86087s	41884	163.81us	130.63us	5.5894ms	maxwell_scudnn_128x
		3.88%	6.51599s	111540	58.418us	24.257us	9.1865ms	void cudnn::bn_fw_t
		3.01%	5.06350s	36660	138.12us	82.243us	9.0224ms	maxwell_sgemv_128x6
		2.93%	4.92540s	364260	13.521us	960ns	2.3991ms	_ZN2at6native29vect
		2.52%	4.23946s	13260	319.72us	153.10us	10.500ms	void cudnn::cnn::wg
		2.48%	4.16063s	27300	152.40us	137.41us	4.3984ms	maxwell_scudnn_128x
		2.39%	4.01173s	77264	51.922us	12.192us	4.3108ms	void cudnn::winogra
		2.30%	3.85825s	36660	105.24us	9.3440us	4.7192ms	void cudnn::winogra
		2.14%	3.59595s	364260	9.8710us	768ns	4.4021ms	_ZN2at6native29vect
		1.46%	2.45122s	4680	523.77us	465.24us	619.26us	void cudnn::cnn::wg
		1.36%	2.27590s	364260	6.2480us	896ns	2.8016ms	_ZN2at6native29vect
		1.26%	2.11418s	12362	171.02us	84.004us	12.024ms	maxwell_scudnn_wino
		1.25%	2.10078s	7866	267.07us	133.64us	4.8152ms	void precomputed_co
		1.20%	2.02146s	364260	5.5490us	736ns	1.0458ms	void at::native::ve
		1.09%	1.82651s	5460	334.53us	238.00us	621.82us	void cudnn::cnn::wg
		0.96%	1.61851s	8580	188.64us	82.148us	5.0181ms	void cudnn::bn_bw_1
		0.95%	1.59084s	6240	254.94us	162.41us	10.394ms	maxwell_scudnn_128x
		0.87%	1.46213s	8792	166.30us	84.163us	7.2406ms	maxwell_scudnn_128x
		0.82%	1.38390s	366287	3.7780us	608ns	2.1559ms	void at::native::ve
		0.73%	1.21785s	3896	312.59us	222.22us	4.7840ms	void precomputed_co
		0.67%	1.12980s	8580	131.68us	58.371us	2.2757ms	void cudnn::bn_fw_t
		0.67%	1.11785s	36660	30.492us	13.249us	2.0194ms	void cudnn::winogra
		0.63%	1.06592s	30380	35.086us	4.4480us	5.5235ms	void cudnn::bn_fw_1
		0.63%	1.05865s	36660	28.877us	11.457us	2.3371ms	void cudnn::winogra
		0.61%	1.03286s	5268	196.06us	99.076us	5.1863ms	maxwell_scudnn_wino
		0.55%	924.02ms	5460	169.23us	157.22us	4.2721ms	maxwell_scudnn_128x
		0.49%	819.03ms	1560	525.02us	434.16us	724.67us	void cudnn::detail:
		0.42%	699.09ms	2340	298.76us	250.95us	4.2314ms	void cudnn::detail:
		0.37%	623.92ms	978	637.95us	242.00us	11.460ms	void explicit_conv
		0.37%	620.00ms	1560	397.44us	355.54us	462.61us	maxwell_scudnn_128x
		0.32%	539.48ms	780	691.64us	674.24us	731.81us	maxwell_scudnn_128x
		0.30%	511.08ms	97092	5.2630us	1.9200us	8.3843ms	task cudnn::compute
		0.30%	498.11ms	1952	255.18us	123.46us	4.8029ms	maxwell_scudnn_wino
		0.29%	493.52ms	1948	253.35us	230.89us	1.5981ms	maxwell_scudnn_128x
		0.27%	458.24ms	780	587.49us	564.44us	680.09us	void at::native::G
		0.24%	396.89ms	42900	9.2510us	2.6560us	631.07us	void cudnn::ops::S
		0.21%	353.43ms	1560	226.56us	218.09us	239.50us	maxwell_scudnn_wino
		0.21%	351.23ms	1952	179.94us	37.218us	1.1780ms	maxwell_scudnn_128x
		0.20%	331.22ms	780	424.64us	418.51us	451.28us	void cudnn::bn_bw_1
		0.19%	315.51ms	976	323.26us	172.58us	748.32us	maxwell_scudnn_128x
		0.18%	299.01ms	976	306.36us	161.90us	624.06us	maxwell_scudnn_128x
		0.15%	252.80ms	780	324.11us	316.72us	387.54us	maxwell_scudnn_128x
		0.15%	249.31ms	120900	2.0620us	1.3120us	131.91us	void at::native::ve
		0.14%	236.76ms	780	303.54us	297.10us	313.52us	void cudnn::bn_fw_t
		0.12%	204.61ms	2888	70.846us	672ns	869.67us	[CUDA memcpy HtoD]
		0.12%	198.07ms	780	253.94us	249.04us	303.21us	maxwell_scudnn_128x
		0.11%	190.44ms	47898	3.9750us	1.8240us	579.93us	void cudnn::cnn::ke

# Nsight

- The recommended system for profiling with a lot more bells and whistles
  - Command line utility similar to nvprof: `nsys`
- `nsys profile --trace=cuda,cudnn,cublas,osrt,nvtx python simply_resnet.py /raid/erik/datasets/imagenet_subset/ --device cuda:4`
- Doesn't require privileged access. Produces reports by default which can be analyzed in GUI tool



# py-spy

- Profile the python-part of code
- Install with pip: `pip install py-spy`
- Can produce multiple different outputs: flame-graphs, speedscope and raw outputs

# py-spy

- `py-spy record --output py-spy-profile.svg -- python simply_resnet.py /data/datasets/imagenet_subset/ --device cuda:0 --pin-memory --num-workers 12`
- By default this creates a *flamegraph*, here written to `py-spy-profile.svg`
- Can also produce *speedscope* files and raw data dumps

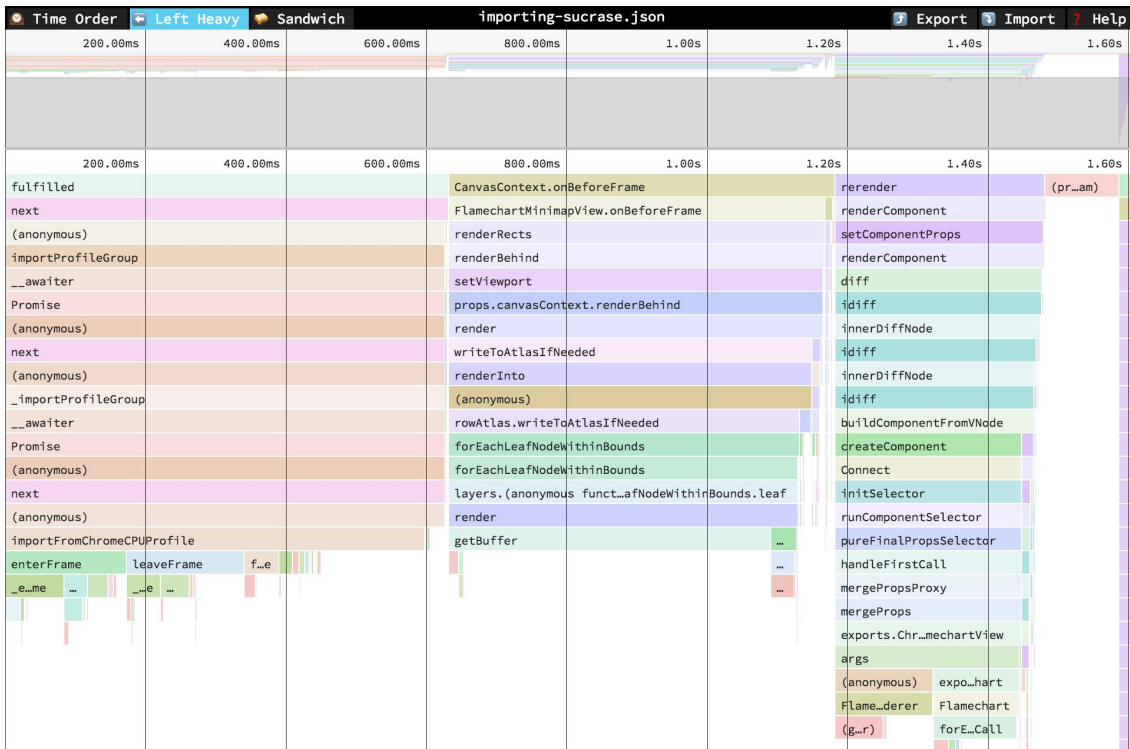
# Flamegraph

- Gives a quick overview of what function has been most active
- X-axis shows percentage of use at a certain stack depth (order is typically alphabetical, not time)
- Y-axis shows call stack depth, so a “flame” is essentially a branch of the call tree



# Speedscope

- Alternative to flame graph, interactive with more ways to view profiles
- View profiles with web app:  
<https://www.speedscope.app/>
- Install locally from  
<https://github.com/jlfwong/speedscope>



# Line\_profiler

- Tool to do fine grained and targeted profiling
- Needs to modify the script to profile
- Install with pip:
  - `pip install line_profiler`
- The first hit on Google is the deprecated repo, this is the correct one: [https://github.com/pyutils/line\\_profiler](https://github.com/pyutils/line_profiler)



# Modifying script for line\_profiler

- The line profiler program looks for functions decorated with the `@profile` decorator
- Add this to the functions you wish to line-profile
- To profile you script, call `kernprof -l` instead of `python`
  - `kernprof -l script.py`
- This dumps profiling information to a file as the script, with an `.lprof` extension

# Inspect line\_profiler dumps

- use the line\_profiler module to read the .lprof files produced by kernprof -l
  - python -m line\_profiler script.py.lprof

```

35 123      1      50763.0  50763.0      0.0      {if not torch.cuda.is_available():
36 124      1      50763.0  50763.0      0.0      raise RuntimeError(f"CUDA not available and device set to {args.device}")
37 125      1      50763.0  50763.0      0.0      else:
38 126      1      252.0    252.0      0.0          device = torch.device(args.device)
39 127      1      12.0     12.0      0.0          torch.backends.cudnn.benchmark = True
40 128      1      36.0     36.0      0.0      print(f"Device is set to {device}")
41 129      1      40725.0  40725.0      0.0      train_set, dev_set, test_set = make_datasets(args.images, rng=rng)
42 130      1      40725.0  40725.0      0.0      batch_size = 4
43 131      1      1.0      1.0      0.0      max_epochs = 1
44 132      1      1.0      1.0      0.0      dataloader_kwargs = dict()
45 133      1      1.0      1.0      0.0      if args.pin_memory:
46 134      1      1.0      1.0      0.0          dataloader_kwargs['pin_memory'] = True
47 135      1      1.0      1.0      0.0      if args.num_workers:
48 136      1      1.0      1.0      0.0          dataloader_kwargs['num_workers'] = args.num_workers
49 137      1      1.0      1.0      0.0      if args.prefetch_factor:
50 138      1      1.0      1.0      0.0          dataloader_kwargs['prefetch_factor'] = args.num_workers
51 139      1      68.0     68.0      0.0      training_loader = DataLoader(train_set, batch_size=batch_size, **dataloader_kwargs)
52 140      1      39.0     39.0      0.0      dev_loader = DataLoader(dev_set, batch_size=batch_size, **dataloader_kwargs)
53 141      1      37.0     37.0      0.0      test_loader = DataLoader(test_set, batch_size=batch_size, **dataloader_kwargs)
54 142      1      657171.0  657171.0      0.3      model = resnet152(pretrained=False, num_classes=train_set.num_classes)
55 143      1      657171.0  657171.0      0.3      #model = resnet18(pretrained=False, num_classes=train_set.num_classes)
56 144      1      657171.0  657171.0      0.3      #model = LogisticRegression(train_set[0][0].shape, train_set.num_classes)
57 145      1      657171.0  657171.0      0.3      #model = alexnet(pretrained=False, num_classes=train_set.num_classes)
58 146      1      657171.0  657171.0      0.3      model.to(device)
59 147      1      657171.0  657171.0      0.3      loss_fn = nn.CrossEntropyLoss()
60 148      1      9452.0    9452.0      0.0      optimizer = Adam(model.parameters(), lr=1e-3, weight_decay=3e-7)
61 149      1      9452.0    9452.0      0.0      for epoch in range(max_epochs):
62 150      2      7.0      3.5      0.0          training_losses = []
63 151      2      7.0      3.5      0.0          for x, y in tqdm(training_loader, desc='training progress'):
64 152      2      2.0      2.0      0.0              optimizer.zero_grad()
65 153      2      2.0      2.0      0.0              prediction = model(x.to(device))
66 154      2      12.4     12.4      0.0              loss = loss_fn(prediction, y.to(device))
67 155      2      8.2      8.2      0.0              loss.backward()
68 156      2      15.1     15.1      0.0              optimizer.step()
69 157      2      4.2      4.2      0.0              training_losses.append(loss.item())
70 158      2      4.2      4.2      0.0          print(f'Training Loss: {np.mean(training_losses)}')
71 159      2      4.2      4.2      0.0          val_losses = []
72 160      2      4.2      4.2      0.0          model.eval()
73 161      2      4.2      4.2      0.0          with torch.no_grad():
74 162      2      4.2      4.2      0.0              for x, y in dev_loader:
75 163      2      4.2      4.2      0.0                  prediction = model(x.to(device))
76 164      2      4.2      4.2      0.0                  loss = loss_fn(prediction, y.to(device))
77 165      2      4.2      4.2      0.0
78 166      2      4.2      4.2      0.0
79 167      2      4.2      4.2      0.0
80 168      2      4.2      4.2      0.0
81 169      2      4.2      4.2      0.0
82 170      2      4.2      4.2      0.0
83 171      2      4.2      4.2      0.0
84 172      2      4.2      4.2      0.0

```

# line\_profiler tip

- The @profile decorator gets defined in the global scope of your script *when* you invoke it using kernprof. It will not be available otherwise
- To quickly go between supporting line\_profiler and running normally, add this snippet to your script
  - It defines a dummy @profile decorator if there isn't one

```
try:
    @profile
    def foo():
        pass
    del foo
except
NameError:
    def
profile(f):
    return f
```

# QnA on profiling

# Efficient PyTorch with DeepSpeed

# Data loaders

- Data loaders are essential to good performance
- PyTorch `torch.util.data.DataLoader` has two important settings
  - `num_workers`, setting this to a positive integer handles data loading in separate processes. Set this to as many CPU threads as you have. Each process will load and collate one batch, making sure that the data loading does not block the main training loop. This hides latency of processing even if it's relatively costly.
  - `pin_memory`, set this to `True` to make the returned tensors be placed in main memory pinned by CUDA, which makes transfer to the GPU faster

# DataParallel vs DistributedDataParallel vs Deepspeed

- The recommended way of doing data parallel workloads in pytorch is the Distributed Data Parallel framework
- There is the more easy to use DataParallel, but this has performance issues since a single process handles all synchronization across GPUs (see this excellent write-up: <https://www.telesens.co/2019/04/04/distributed-data-parallel-training-using-pytorch-on-aws/>)
- Distributed Data Parallel is relatively easy to use, but documentation is a bit scattered
- There's another framework which does the same thing, and a lot more: Deepspeed from Microsoft, <https://github.com/microsoft/DeepSpeed>, [deepspeed.ai](https://deepspeed.ai)

# Deepspeed

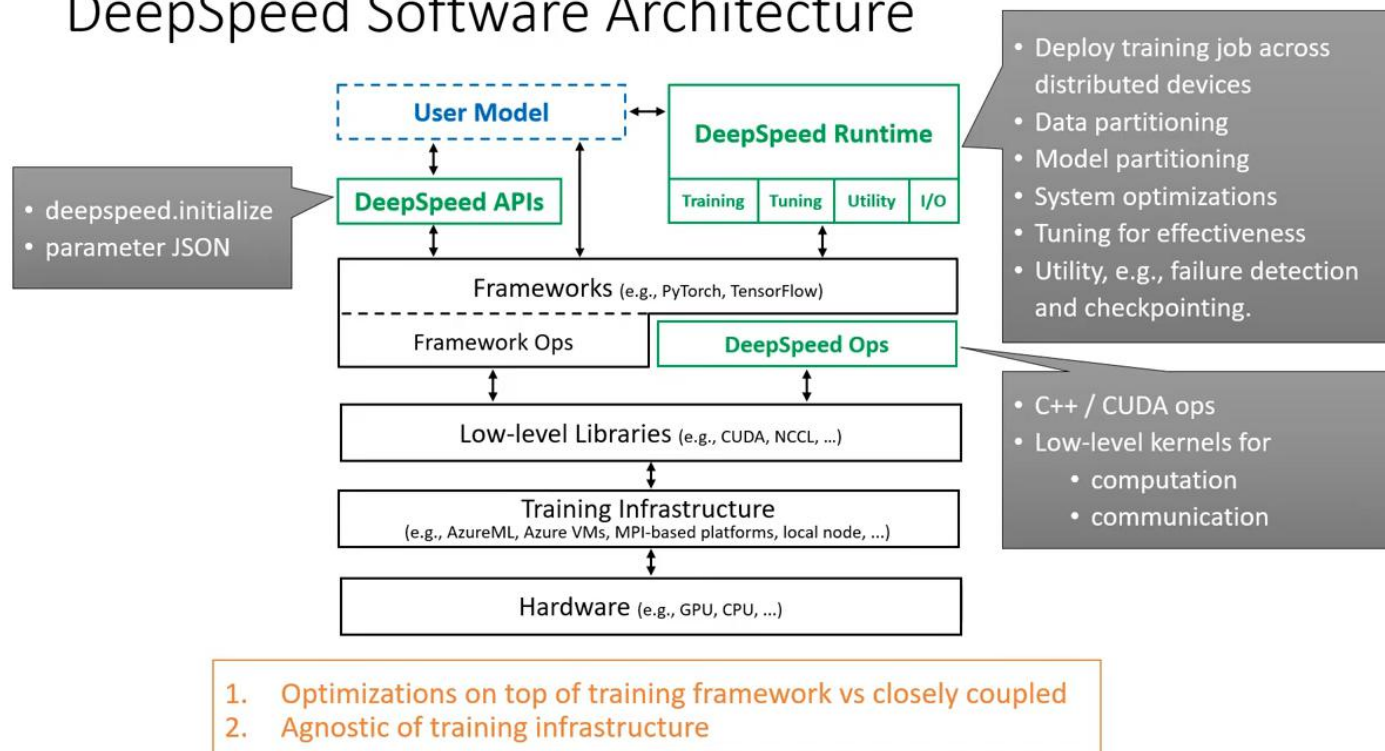
- Relatively new framework (there are definitely rough edges)
- Designed to take away as much of the hustle of writing efficient deep learning programs as possible
- Distributed learning, same code works as well for single-node+single-gpu/single-node+multiple-gpu/multiple-nodes+multiple-gpus
- Efficient ways of synchronizing gradients with custom optimizers
- Built in support for fp16 and mixed precision training (AMP) using NVIDIA's Apex package as well as automatically handling of loss scaling and correct gradient clipping
- Ways of easily doing model parallel training with pipelining, and combining it with data parallel training



# Deepspeed

- Optimization framework Zero, for sharding optimizer state/gradients/parameters over multiple GPUs.
  - Can drastically reduce memory usage with minor impact on performance
  - The first stage of this optimization shards the optimizer states (e.g. ADAMs moving averages) so that each worker only has parts of the state
- One driver for deepspeed are huge Language Models, so framework also contains custom efficient kernels for Transformers
- Also has some support for one-cycle style LR policies, including learning rate range finders. There seems to be gathering evidence that this is a good idea for large batch training.

# DeepSpeed Software Architecture



[Hands-on Tutorials] DeepSpeed 01, KDD2020  
<https://youtu.be/CaseqC45DNc>

# Deepspeed gotchas

- Includes many custom PyTorch extensions (e.g. optimizers) which will be JIT-compiled when running
  - Needs full build environment (so full CUDA-install and C++ compiler and standard libraries)
  - This shouldn't be a problem at AIDA where you have full control over your VM

# Deepspeed install

- First install CUDA of desired version
- Install pytorch (recommended through anaconda) to match the installed cuda version
- Install deepspeed with pip:
  - `pip install deepspeed`

# Modifying a script to use deepspeed

- The script will be launched through deepspeeds launcher script. Certain command line arguments are assumed to be available so run the following on your `argparse.ArgumentParser` object:

- `parser = deepspeed.add_config_arguments(parser)`

- I had to add the following explicitly:

- `parser.add_argument('--local_rank', type=int, default=-1, help='local rank passed from distributed launcher')`

# Deepspeed configuration

- Deepspeed is configured using a JSON file with a particular set of values
- Most parameters for training is changed here, such as batch size, gradient accumulation steps, optimizer, scheduler, 16bit precision training, model parallelism and Zero configuration

```
{
  "train_batch_size": 4,
  "steps_per_print": 2000,
  "optimizer": {
    "type": "Adam",
    "params": {
      "lr": 0.001,
      "betas": [
        0.8,
        0.999
      ],
      "eps": 1e-8,
      "weight_decay": 3e-7
    }
  },
  "scheduler": {
    "type": "WarmupLR",
    "params": {
      "warmup_min_lr": 0,
      "warmup_max_lr": 0.001,
      "warmup_num_steps": 1000
    }
  },
  "wall_clock_breakdown": false
}
```

# Modifying a script to use deepspeed

- Deepspeed provides a wrapper around any torch.nn.Module object
- To use deepspeed, wrap your model in the deepspeed.initialize function:
  - ```
model_engine, optimizer, __, __ =  
    deepspeed.initialize(args=args, model=model,  
    model_parameters=model.parameters())
```
- The args is your command line arguments which deepspeed will go through to find its configuration
- The model\_engine object is now our wrapped model and will behave

# Deepspeed dataloaders

- Deepspeed will work with PyTorch's dataloader, but to handle parallel sampling without extra steps use:
  - `training_loader = model_engine.deepspeed_io(train_set)`
- This wrapper also accepts some of the same arguments as PyTorch dataloaders, in particular `collate_fn` and `pin_memory`
  - Other arguments like batch size or sampler can be left to the wrapper



# Modifying script to use deepspeed

- Since we assume data parallel, we need to let deepspeed determine where to place tensors
- When we move a tensor to device, use `model_engine.local_rank`
  - `prediction = model_engine(x.to(model_engine.local_rank))`
- This takes care of placing the tensor on the correct GPU based on the process running the code

# Modifying a script to use deepspeed

- Deepspeed handles the optimizer steps, as well as the backpropagation
- We need to replace our typical calls to the optimizer with specific calls from the model engine
  - Don't call the backward of the loss node (`loss.backward()`), instead use `model_engine.backward(loss)`
  - Don't call `optimizer.zero_grads()` or `optimizer.step()`, instead call `model_engine.step()`. It will take care of zeroing gradients when appropriate.

# Running the new script

- We need to run the script using deepspeeds launcher:
  - `deepspeed simply_resnet_deepspeed.py  
/data/datasets/imagenet_subset/ --deepspeed_config  
ds_config.json`
- This takes care of setting up the runtime environment and inject the correct information in all processes

# Training large batches

- Training with large batches has historically generated worse models than with small batches
- Normalizations (such as BatchNorm) seems to help with this
- Learning rate schedule also seems to be important (LAMB and One-Cycle policy seems to be gaining momentum)

## ABSTRACT

Training large deep neural networks on massive datasets is computationally very challenging. There has been recent surge in interest in using *large batch* stochastic optimization methods to tackle this issue. The most prominent algorithm in this line of research is LARS, which by employing *layerwise adaptive* learning rates trains RESNET on ImageNet in a few minutes. However, LARS performs poorly for attention models like BERT, indicating that its performance gains are *not* consistent across tasks. In this paper, we first study a principled layerwise adaptation strategy to accelerate training of deep neural networks using large mini-batches. Using this strategy, we develop a new layerwise adaptive large batch optimization technique called LAMB; we then provide convergence analysis of LAMB as well as LARS, showing convergence to a stationary point in general nonconvex settings. **Our empirical results demonstrate the superior performance of LAMB across various tasks such as BERT and RESNET-50 training with very little hyperparameter tuning.** In particular, for BERT training, our optimizer enables use of very large batch sizes of 32868 without any degradation of performance. By increasing the batch size to the memory limit of a TPUv3 Pod, BERT training time can be reduced from 3 days to just 76 minutes (Table 1). The LAMB implementation is available online<sup>1</sup>.

You, Yang, et al. "Large batch optimization for deep learning: Training bert in 76 minutes." ICLR 2020

# QnA on deepspeed

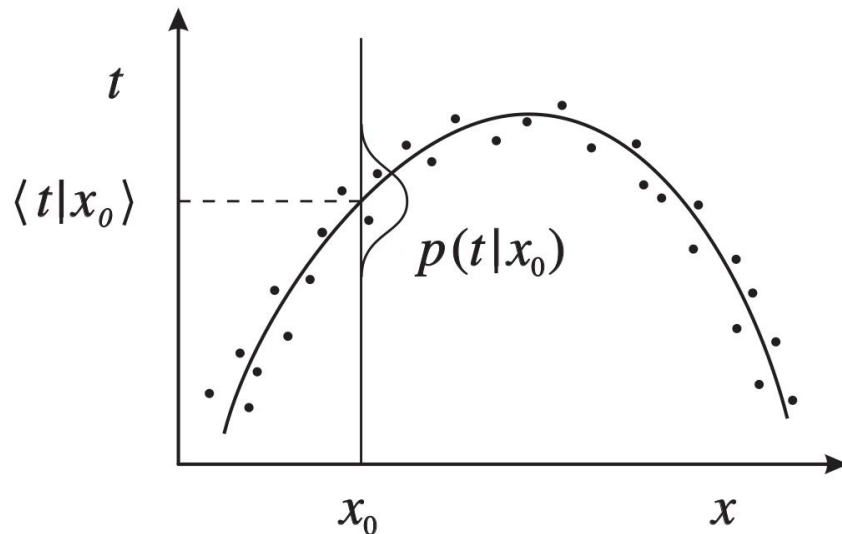
**Erik Ylipää**  
erik.ylipaa@ri.se

# One-hot-encodings

- Never use one-hot-encodings as inputs to neural networks
  - Instead use an Embedding layer with integer encoded categorical values
- There are two issues, 1) inefficiency and 2) optimization:
  1. Multiplying a one-hot-vector with a matrix is mostly wasted computation, multiplying all rows of the matrix except one with 0
  2. If the vocabulary is large, or the frequency distribution is very skewed, optimizers which accumulates gradient statistics (i.e ADAM, RMSProp, SGD+momentum) have no chance of applying sparse updates to those statistics

# Regression targets

- If we model regression with a linear output layer and Mean Squared Error loss the implicit assumption is that the conditional distribution is Gaussian
- In particular, we will predict the conditional mean

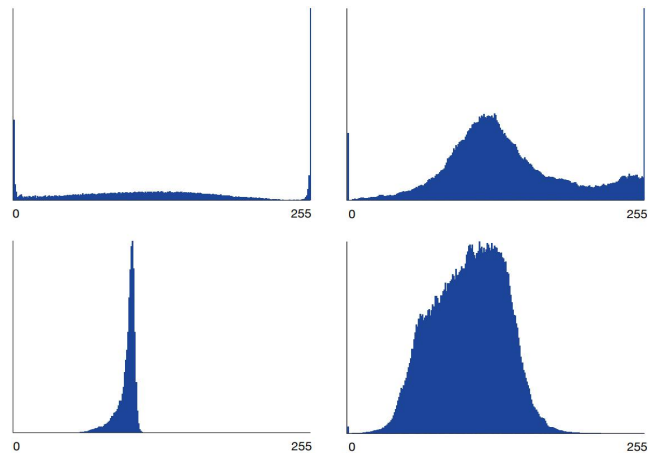


Bishop, Christopher M. 'Mixture density networks.' (1994).



# Non-Gaussian $P(Y|X)$

- Often the conditional distribution is not Gaussian
- If the real distribution is multimodal or skewed, the conditional mean will often be a poor prediction
- A mixture of Gaussian can theoretically model any continuous distribution, but in practice can underperform



Four different conditional distributions on autoregressive tasks

# Discretizing trick

- Model continuous variables as discrete
- Essentially predict histogram bins
- Determining *bins* (hinkar) becomes an issue
- One strategy is to look at the marginal empirical CDF of Y
- How many bins to use? Look at the data, what granularity makes sense?

